

The Code-Constraint Problem in Biological Systems: How Low-Dimensional Interfaces Shape High-Dimensional Dynamics

Ian Todd

Sydney Medical School

University of Sydney

Sydney, NSW, Australia

`itod2305@uni.sydney.edu.au`

December 20, 2025

Abstract

A recurring problem spans biophysics, systems biology, and neuroscience: high-dimensional systems (protein ensembles, gene networks, neural populations) are characterized through low-dimensional descriptions (order parameters, principal components, expression markers), yet the relationship between interface and underlying dynamics remains poorly understood. When is dimensional reduction faithful compression versus systematic distortion? Why do some low-dimensional codes work while others fail?

This review proposes a unifying interpretation: low-dimensional interfaces between coupled systems function as *stabilizing constraints* rather than information channels. We synthesize evidence from dynamical systems theory, statistical mechanics, and information theory to argue that the recurring interpretive problems arise from a common

source—the assumption that dimensional reduction preserves information rather than imposing constraints.

Using coupled oscillator simulations as a minimal physical exemplar, we demonstrate a characteristic signature of constraint: bandwidth-limited coupling induces systematic complexity collapse in responding systems while maintaining bounded tracking. Critically, this requires structured projections that capture coherent collective variables; random projections of the same dimensionality produce the opposite effect (whitening, not collapse). The same signature appears in gene regulatory network dynamics, confirming generality beyond oscillatory systems.

We connect this mechanism to protein folding, single-cell manifolds, morphogenetic fields, and neural coding, and offer practical criteria for researchers: how to choose reaction coordinates, interpret manifold structure, design multi-scale interfaces, and distinguish constraint from information. The framework suggests that biological codes function primarily as dimensional valves enabling persistent organization, rather than as channels optimized for information transmission.

Keywords: Dimensional reduction, coarse-graining, biological codes, constraint dynamics, order parameters, model identifiability

1 A Common Problem Across Fields

Across biophysics, systems biology, and neuroscience, a recurring tension appears: the systems we study operate in high-dimensional state spaces, yet we characterize them through low-dimensional descriptions. This dimensional mismatch is not merely a practical limitation—it generates systematic interpretive problems that recur across domains.

In protein biophysics, conformational ensembles explore thousands of degrees of freedom, yet function is characterized by a handful of order parameters. Reaction coordinates that appear to govern folding kinetics may hide multiple parallel pathways (Pande et al., 2010). Parameter estimation is generically “sloppy”—many parameter combinations pro-

duce indistinguishable low-dimensional outputs (Gutenkunst et al., 2007; Transtrum et al., 2015). The question of when order-parameter descriptions are reliable has been extensively studied, yet remains unresolved (Best and Hummer, 2005; Noé and Clementi, 2017).

In single-cell genomics, RNA sequencing reveals that cell states occupy low-dimensional manifolds embedded in \sim 20,000-dimensional gene expression space (Trapnell, 2015; Weinreb et al., 2018). Yet dimensionality reduction techniques can create spurious clusters, collapse distinct states, and distort neighborhood relationships (Chari and Pachter, 2023; Kobak and Berens, 2019; Cooley et al., 2019). Pseudotime analyses and trajectory inference rely on manifold structure that may reflect technical artifacts rather than biological reality.

In developmental biology, tissues coordinate through field-like gradients—morphogen concentrations, bioelectric potentials, mechanical stresses (Levin, 2021; Turing, 1952). These fields may themselves possess rich, high-dimensional dynamics at molecular and ionic scales, yet individual cells respond to local values through bandwidth-limited interfaces (ion channels, receptors). The mismatch between field complexity and cellular readout is structural, not merely practical.

In neuroscience, neural populations fire in high-dimensional pattern spaces, yet must coordinate across distant brain regions through apparent low-dimensional codes (Gallego et al., 2017; Sadtler et al., 2014). Order parameters that appear to govern dynamics may be shadows of higher-dimensional processes (Flack, 2017; Balduzzi and Tononi, 2008). Learning requires either operating within existing neural manifolds or expanding them—a constraint whose origin remains unclear.

These problems share a common structure: high-dimensional systems coupled through low-dimensional interfaces, where the relationship between interface and underlying dynamics is poorly understood. When can we trust a low-dimensional model? When is dimensional reduction a faithful compression versus a systematic distortion?

1.1 This Review: Codes as Constraints

This review proposes a unifying interpretation of these phenomena, rather than a new biological mechanism. We synthesize insights from dynamical systems theory, statistical mechanics, and information theory to argue that the recurring problems above arise from a common source: the assumption that dimensional reduction preserves information rather than imposing constraints.

The core reframing is that low-dimensional interfaces between coupled systems function as *stabilizing constraints* rather than information channels.

Consider any biological interface: a cell membrane transducing environmental signals, a morphogenetic gradient coordinating tissue development, a neural code coupling sensory input to motor output. In each case, a high-dimensional “driving” system couples to a high-dimensional “responding” system through a low-dimensional bottleneck. The question is: what does this bottleneck *do*?

The standard information-theoretic framing asks how much information about the driving system’s state is preserved or lost. We propose a different question: how does the bottleneck *shape the dynamics* of the responding system?

Using a minimal model of coupled oscillator lattices, we demonstrate that bandwidth-limited coupling produces a distinctive signature:

1. **Complexity collapse:** The responding system’s effective dimensionality decreases systematically with code bandwidth.
2. **Bounded tracking:** Alignment between systems remains stable despite information loss.
3. **Structure dependence:** The effect requires projections onto coherent collective variables; random projections fail.

This pattern—complexity collapse with bounded tracking—is qualitatively different from information loss (which would produce tracking failure) or simple filtering (which would not

systematically reduce effective dimensionality). We argue it represents a general mechanism by which biological coding structures enable persistent organization.

Positioning relative to existing theory: The framework relates to Haken’s synergetics, where fast variables become “slaved” to slow order parameters—but extends this to coupled systems where the order parameter constrains a *separate* responding system. It differs from the information bottleneck principle, which optimizes compression for predictive relevance; we instead measure how compression constrains downstream dynamics, regardless of predictive content. It complements Markov blanket formalism by adding bandwidth as a quantitative knob: not just whether states are separated, but how many degrees of freedom can cross the boundary.

To be explicit about the novel claim: we are not proposing a better compression or a better predictor. We propose a *distinct dynamical regime diagnostic*: complexity collapse in the responding system while tracking error remains bounded, where “bounded” means alignment with the *coarse-grained reconstruction*, not full state reconstruction. This regime appears only for structured projections (coherent collective variables) rather than arbitrary low-dimensional reductions.

2 Biophysical Contexts

Before presenting the formal model, we ground the framework in concrete biological systems where dimensional mismatch is experimentally relevant.

2.1 Protein Conformational Ensembles

Proteins exist as dynamic ensembles exploring rugged energy landscapes (Frauenfelder et al., 1991; Henzler-Wildman and Kern, 2007). The full conformational state space has thousands of dimensions, yet function is often characterized by a handful of order parameters: radius of gyration, native contacts, collective mode amplitudes.

The question of when order-parameter descriptions are reliable has been extensively studied (Best and Hummer, 2005; Noé and Clementi, 2017). Reaction coordinates that appear to govern folding kinetics may hide multiple parallel pathways (Pande et al., 2010). The “folding funnel” picture (Dill and Chan, 1997) is itself a low-dimensional projection that may obscure high-dimensional heterogeneity.

Our framework suggests a reinterpretation: rather than asking whether order parameters *represent* the ensemble, we ask whether they *constrain* accessible trajectories. A well-chosen collective variable doesn’t merely summarize the ensemble—it restricts what conformational dynamics remain possible. This is the difference between a thermometer (which measures temperature without affecting it) and a thermostat (which constrains the system to a temperature range).

Translating the model: In protein-ligand binding, System *A* is the ligand’s conformational dynamics in solution (high-dimensional, exploring many states). System *B* is the binding-pocket dynamics of the target protein. The “code” is the subset of ligand conformations that can physically dock—perhaps characterized by 2–5 collective variables describing complementary surface geometry. When ligand binds, these few degrees of freedom constrain the pocket’s accessible conformations, inducing complexity collapse in the binding site while maintaining the lock-and-key alignment. The signature would be: (1) reduced conformational entropy of the bound pocket relative to apo form, (2) stable binding geometry despite ligand fluctuations, and (3) this constraint failing for decoy ligands that match steric volume but not the structured collective variables of true binding. One could test this by comparing crystallographic B-factors (conformational flexibility) of binding pockets with cognate versus non-cognate ligands: the framework predicts that cognate binding induces greater complexity collapse.

2.2 Single-Cell State Manifolds

Single-cell RNA sequencing reveals that cell states occupy low-dimensional manifolds embedded in high-dimensional gene expression space (Trapnell, 2015; Weinreb et al., 2018). Pseudotime analyses, trajectory inference, and RNA velocity all rely on this manifold structure.

Yet recent critiques highlight systematic artifacts: t-SNE and UMAP can create spurious clusters, apparent trajectories may reflect technical noise, and neighborhood relationships are often distorted (Chari and Pachter, 2023; Cooley et al., 2019). The fundamental issue is that \sim 20,000-dimensional expression space is projected onto 2–3 visualization dimensions, a compression ratio that inevitably loses information.

The constraint perspective offers a different question: which expression programs *constrain* cell fate transitions? A gene regulatory motif that restricts accessible trajectories (like a developmental “valve”) functions differently from one that merely correlates with cell state. The former shapes dynamics; the latter reflects them.

2.3 Morphogenetic Fields and Bioelectric Gradients

Developmental biology increasingly recognizes that tissues coordinate through field-like gradients—morphogen concentrations, bioelectric potentials, mechanical stresses (Levin, 2021; Turing, 1952). These fields may themselves possess rich, high-dimensional dynamics at molecular and ionic scales. However, the *interface* through which individual cells couple to tissue-scale organization is bandwidth-limited: cells respond to local voltage or concentration values, not to the full spatiotemporal complexity of the field.

Michael Levin’s work on bioelectric control of morphogenesis demonstrates that manipulating membrane voltage patterns can reprogram tissue-scale outcomes (Levin, 2021, 2014). The key insight is not that bioelectric fields are simple, but that cellular responses are constrained to low-dimensional projections of field dynamics. Cells “listen” through ion channels that act as bandpass filters, extracting only certain features of the voltage landscape.

This is precisely the mechanism we formalize below: a high-dimensional responding system (individual cell dynamics) couples to a driving system (the broader tissue/environment) through a bandwidth-limited interface, and the interface’s bandwidth—not the intrinsic complexity of either system—constrains the responding system’s behavioral repertoire.

2.4 Neural Population Codes

Neuroscience faces the dimensional mismatch problem acutely. Neural populations fire in high-dimensional pattern spaces, yet must coordinate across distant brain regions. How does the brain achieve this coupling?

A key insight emerges from oscillatory dynamics: low-frequency oscillations can coordinate activity across spatially distributed populations, while high-frequency (gamma) activity carries precise local information but couples only locally (Miller et al., 2018; Buzsáki and Draguhn, 2004). Counterintuitively, this means low-frequency coupling may *enable* high-dimensional coordination—the slow oscillation provides a shared temporal reference that allows distant populations to integrate. The “bandwidth” of inter-regional coupling is set by oscillatory frequency, not by the intrinsic complexity of local circuits.

Population coding theories typically frame neural activity as *representing* task variables (Georgopoulos et al., 1986; Pouget et al., 2000). An alternative framing asks which aspects of population activity *constrain* downstream processing. Recent work on neural manifolds shows that population activity is confined to low-dimensional subspaces (Gallego et al., 2017; Sadtler et al., 2014). But these manifolds may reflect the bandwidth limitations of inter-regional coupling rather than intrinsic computational simplicity. Learning a new behavior requires either operating within the existing manifold or expanding it—a constraint imposed by the coupling interface, not just a summary of neural activity.

3 A Minimal Physical Exemplar

To make the constraint interpretation concrete, we demonstrate its signature in a minimal model: two coupled lattices of phase oscillators where interaction is restricted to a bandwidth-limited projection. This is not intended as a model of any specific biological system, but as a *didactic exemplar*—a physically realizable system that exhibits the constraint mechanism in a form amenable to precise measurement. The same qualitative signature emerges in gene regulatory network dynamics (Section 7.2), confirming that the mechanism is not specific to oscillatory dynamics.

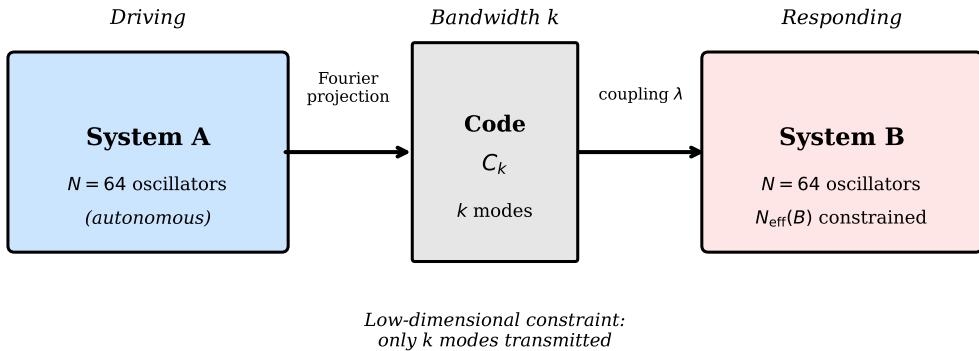


Figure 1: Model architecture. System A (driving) couples to System B (responding) only through a bandwidth-limited code C_k consisting of k Fourier modes. The code constrains B 's dynamics without transmitting full state information.

3.1 Model Architecture

Consider two systems, each a one-dimensional lattice of N locally coupled phase oscillators. System A (the “driving” system) evolves autonomously:

$$\dot{\theta}_i^A = \omega_i + K \sum_{j \in \mathcal{N}(i)} \sin(\theta_j^A - \theta_i^A) + \eta_i^A(t), \quad (1)$$

where ω_i are intrinsic frequencies, K is coupling strength, $\mathcal{N}(i)$ denotes nearest neighbors, and $\eta_i^A(t)$ is Gaussian noise.

System B (the “responding” system) couples to A only through a low-dimensional code:

$$\dot{\theta}_i^B = \omega_i + K \sum_{j \in \mathcal{N}(i)} \sin(\theta_j^B - \theta_i^B) + \lambda \sin(\hat{\theta}_i^A - \theta_i^B) + \eta_i^B(t), \quad (2)$$

where $\hat{\theta}^A$ is a bandwidth-limited reconstruction of A ’s state.

The code C_k consists of the first k Fourier modes of A ’s phase field:

$$C_m = \frac{1}{N} \sum_{i=1}^N e^{i\theta_i^A} e^{-i2\pi mi/N}, \quad m = 0, \dots, k. \quad (3)$$

The bandwidth-limited reconstruction is obtained by inverting with only the retained modes:

$$\hat{\theta}_i^A = \arg \left(\sum_{m=0}^k C_m e^{i2\pi mi/N} \right). \quad (4)$$

We define bandwidth k as the maximum retained mode index (including DC), so the code contains $k + 1$ complex coefficients.

This architecture (Figure 1) captures the essential structure: two high-dimensional systems coupled through a low-dimensional interface. The bandwidth k is a tunable parameter controlling code dimensionality.

3.2 Quantifying Complexity and Tracking

We measure *effective dimensionality* via spectral entropy of the Fourier amplitude spectrum. Specifically, the amplitude $|C_m|$ of each Fourier mode is computed, normalized to form a probability distribution $p_m = |C_m| / \sum_m |C_m|$, and the Shannon entropy $H = -\sum_m p_m \log p_m$ is calculated. The effective dimensionality is then $N_{\text{eff}} = e^H$, which gives the *effective support size*—the number of modes that contribute substantially to the dynamics, weighted by their relative amplitudes. This is not a topological dimension but a complexity measure analogous to the participation ratio in localization theory. The same qualitative results obtain using alternative metrics—PCA participation ratio (covariance-based), spatial gradient energy, and

Kuramoto order parameter—none of which depend on the Fourier basis (Supplementary Figure S3). This confirms that complexity collapse is not an artifact of measuring complexity in the same basis as the code.

Tracking quality is measured as mean circular distance between systems: $\Delta = N^{-1} \sum_i |\sin((\theta_i^A - \theta_i^B)/2)|$.

4 Results: The Constraint Signature

4.1 Complexity Collapse Under Bandwidth Reduction

Figure 2 shows the central result. As code bandwidth k decreases:

- System B 's effective dimensionality decreases systematically ($N_{\text{eff}} = 16.7 \rightarrow 10.9$)
- System A 's complexity remains constant ($N_{\text{eff}} \approx 17.0$)
- The relationship is approximately monotonic

The bottleneck constrains only the responding system; the driving system's dynamics are unaffected by how it is observed. This asymmetry is the hallmark of constraint rather than mutual information loss.

4.2 Bounded Tracking Despite Complexity Collapse

Mismatch between systems increases only modestly as bandwidth decreases (Figure 3): $\Delta = 0.38 \rightarrow 0.50$, a $\sim 32\%$ change, compared to $\sim 35\%$ reduction in $N_{\text{eff}}(B)$.

This asymmetry is crucial. If the code merely lost information, we would expect commensurate tracking failure. Instead, B tracks *the code* faithfully—it aligns with A 's coarse-grained representation. The code constrains B 's behavioral complexity without compromising alignment.

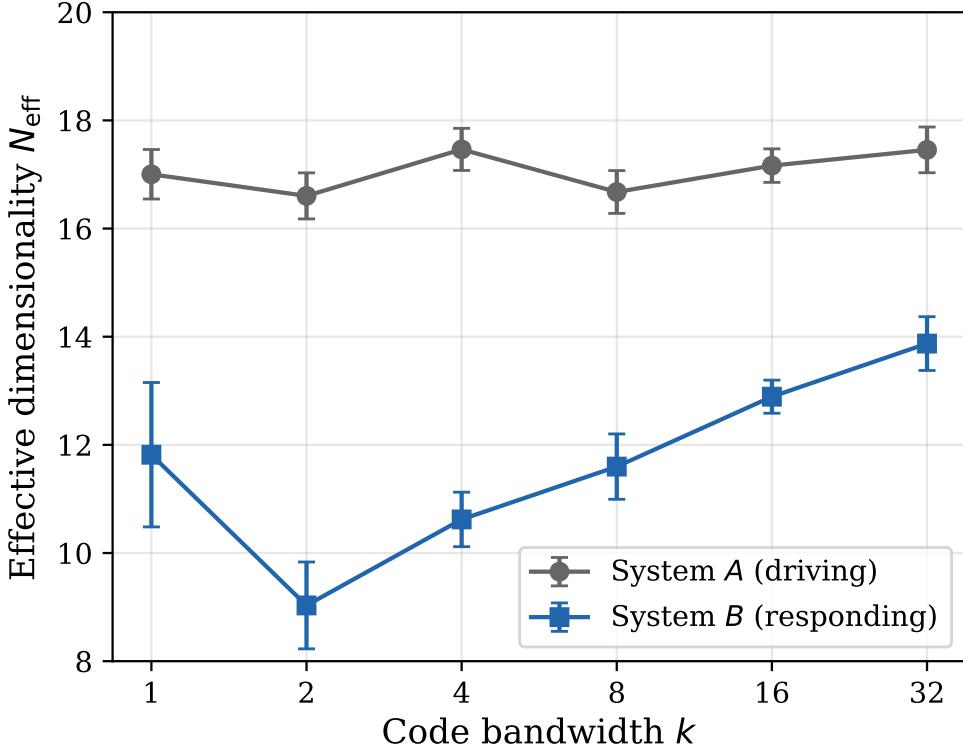


Figure 2: Effective dimensionality as a function of code bandwidth. System A (gray) maintains constant complexity; system B (blue) exhibits systematic collapse. Error bars: standard error over 15 trials.

4.3 Structure Matters: Random Projections Fail

A critical control: we repeated experiments using random k -mode projections instead of the lowest k Fourier modes. Results differ strikingly (Figure 4):

- Random projections produce *no* complexity collapse ($N_{\text{eff}} \approx 15\text{--}17$)
- Mismatch is substantially higher
- At low k , random codes increase B 's complexity above A 's

The constraining effect is not about dimensionality per se, but about *structure*. Low-frequency modes capture spatially coherent patterns; random modes mix signal with noise. Biological codes, we argue, must similarly capture coherent collective variables rather than arbitrary dimensional reductions.

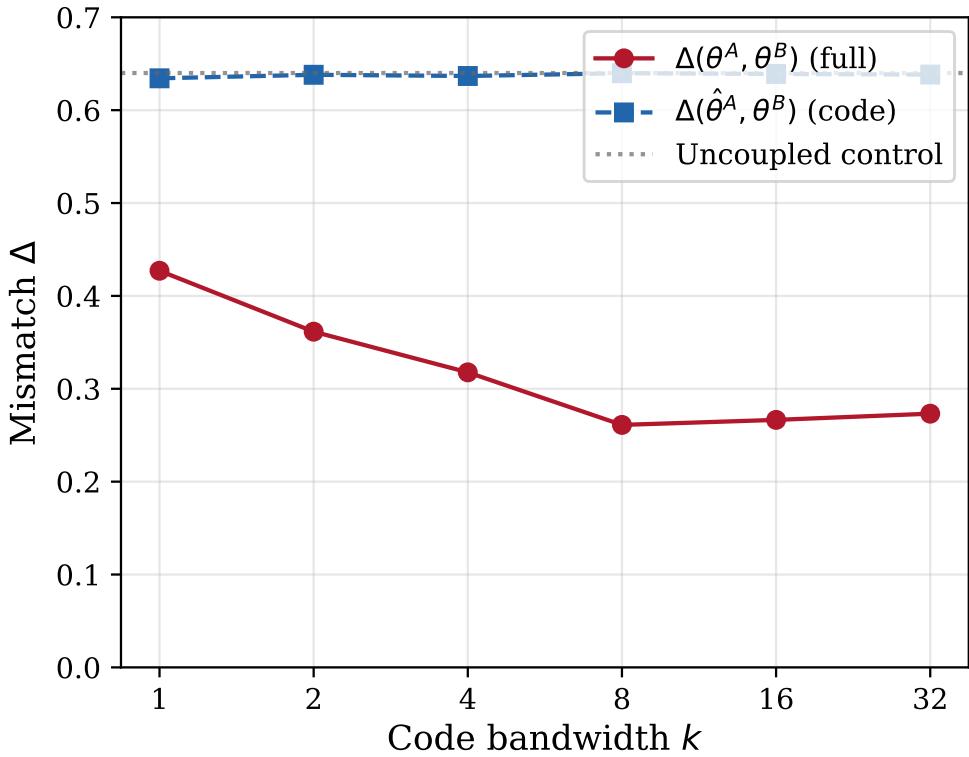


Figure 3: Mismatch between systems as a function of code bandwidth. Coupled systems (solid) maintain lower mismatch than uncoupled controls (dashed) across all bandwidths.

Notably, at low k the random projection actually *increases* B 's complexity above baseline ($N_{\text{eff}} \approx 18$ vs. 17 for the uncoupled system). This occurs because sparse random projections destroy local spatial correlations in the driving signal, effectively acting as a whitening filter that forces the responding lattice to desynchronize locally while still responding to the global coupling. The random code injects high-frequency structure rather than removing it.

Robustness: The qualitative pattern—complexity collapse for structured projections, not for random ones—persists across the parameter ranges tested: noise amplitude $\sigma \in [0.1, 0.5]$, coupling strength $\lambda \in [0.25, 4.0]$, and lattice sizes $N \in [32, 128]$. The effect is not an artifact of specific parameter choices. We expect other structured bases (Laplacian eigenmodes, wavelets, principal components of the driving system) to produce similar constraint signatures, though the specific relationship between k and N_{eff} will depend on the basis.

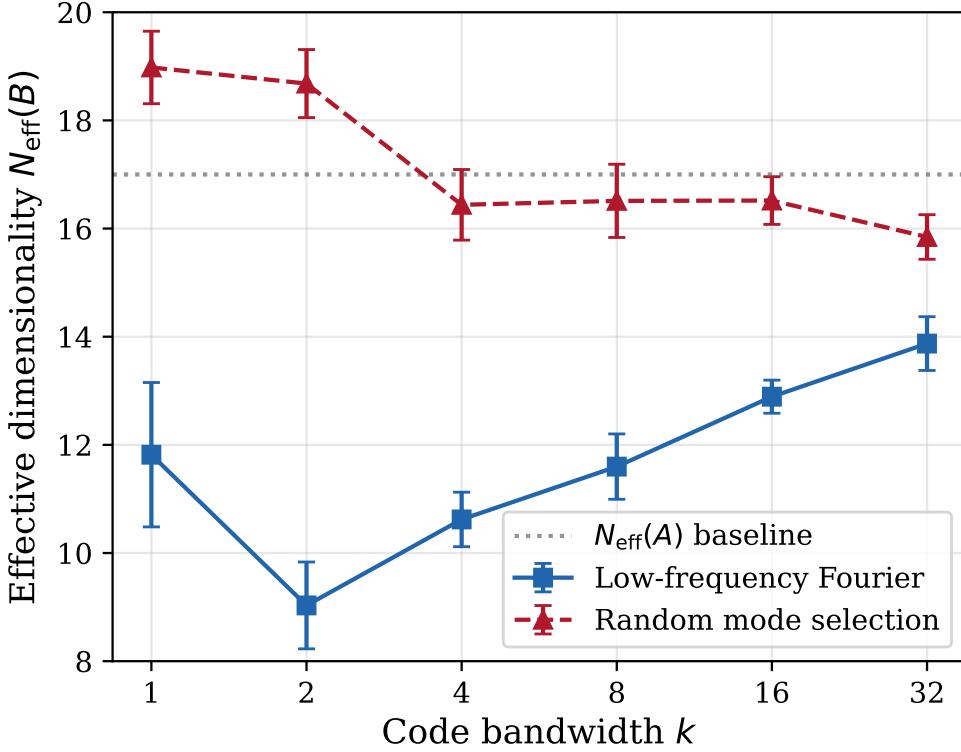


Figure 4: Low-frequency Fourier coupling (blue) vs. random mode selection (red). Random projections of the same dimensionality fail to induce complexity collapse.

5 Implications for Biological Coding Structures

5.1 Codes as Dimensional Valves

The simulation results suggest a general principle: biological interfaces that couple high-dimensional systems through low-dimensional projections function as *dimensional valves*—they constrain the behavioral complexity of responding systems without requiring full state reconstruction.

This reframes the function of biological coding structures:

- **Gene regulatory motifs** don’t merely “represent” cellular state; they constrain accessible developmental trajectories.
- **Morphogenetic gradients** don’t merely “signal” position; they restrict the dynamical complexity of responding tissues.

- **Neural codes** don’t merely “encode” stimuli; they constrain downstream processing to behaviorally relevant dimensions.

The key insight is that constraint and representation are distinct functions. A code can effectively constrain dynamics while having poor predictive power—indeed, discarding destabilizing microstate detail may be precisely why the code works.

5.2 When Coarse-Grained Models Are Reliable

The framework offers practical criteria for evaluating low-dimensional models:

1. **Does the projection capture coherent collective variables?** Random or arbitrary projections will not produce stable constraint. Effective codes must select dynamically relevant degrees of freedom.
2. **Is there a responding system being constrained?** If the “code” is merely an observer’s summary with no causal role, it may misrepresent the underlying dynamics. Constraint requires causal coupling.
3. **Does complexity collapse predict stability?** If reducing the code’s bandwidth produces chaotic rather than simplified dynamics in the responding system, the projection is not capturing constraining structure.

These criteria distinguish cases where order-parameter descriptions are reliable (because they capture genuine constraints) from cases where they are misleading (because they are arbitrary projections onto a richer dynamics).

5.3 The Membrane as Paradigm Case

The cell membrane exemplifies this framework. It is a physical dimensional bottleneck: the external environment and internal cytoplasm are each enormously high-dimensional,

yet coupling occurs through a restricted set of interface variables—receptor occupancy, ion channel conductance, transmembrane voltage.

The membrane does not “represent” the environment; it *constraints* which environmental features can influence cellular dynamics. Evolution has tuned both the bandwidth (interface dimensionality) and structure (which degrees of freedom couple) to maximize actionable constraint per unit metabolic cost.

This perspective suggests that “information processing” metaphors may be misleading for membrane biology. The membrane is not a channel that transmits environmental information; it is a valve that restricts cellular dynamics to viable trajectories given environmental context.

6 Relation to Existing Theory

Our framework connects to several established theoretical traditions:

Synergetics and slaving: Haken’s synergetics describes how fast variables become “slaved” to slow order parameters (Haken, 1983). We extend this to coupled systems: the code is a bandwidth-limited order parameter that constrains a separate responding system.

Information bottleneck: The information bottleneck principle (Tishby et al., 1999) optimizes compression for predictive relevance. Our framework differs: we measure how compression constrains downstream dynamics, not how much predictive information is preserved.

Markov blankets: The Markov blanket formalism (Friston, 2013) identifies boundaries separating internal from external states. We add a quantitative dimension: bandwidth k specifies how many degrees of freedom can cross the blanket.

Computational mechanics: ε -machines characterize minimal sufficient statistics for prediction (Crutchfield and Young, 1989). We propose a sibling concept: *minimal stabilizing interfaces*—projections that constrain dynamics while potentially being poor predictors.

7 Limitations and Future Directions

We emphasize what the simulation demonstrates directly: complexity collapse and bounded tracking under bandwidth-limited coupling. We hypothesize, but do not directly test, that dimensionality reduction is a prerequisite for attractor formation and persistent organization. Perturbation studies would be needed to verify resilience of the constrained dynamics.

The Fourier basis is convenient analytically but not uniquely privileged. The key property is capturing spatially coherent structure; other structured bases (wavelets, Laplacian eigenfunctions) might serve equally well. The critical distinction is between coherent projections that enable constraint and incoherent projections that do not.

Future work should test these predictions in experimental systems: does manipulating interface bandwidth (e.g., via optogenetic control of receptor expression) produce the predicted complexity collapse in responding systems?

8 Generality Across Scales and Dynamics

To test the generality of the constraint signature, we extended the simulations in two directions: larger system sizes and alternative dynamics.

8.1 Large-Scale Kuramoto Systems

We repeated the main analysis at $N = 512$ and $N = 1024$ oscillators (Supplementary Figure S1). The complexity collapse signature scales appropriately:

- At $N = 512$: $N_{\text{eff}}(B)$ ranges from ≈ 102 at $k = 4$ to ≈ 143 at $k = 128$
- At $N = 1024$: the complexity ratio $N_{\text{eff}}(B)/N_{\text{eff}}(A)$ ranges from 0.64 to 0.95 as k increases
- Random projections continue to show no collapse, with $N_{\text{eff}}(B)$ remaining near the uncoupled baseline

The constraint signature is not a finite-size artifact; it persists as system dimensionality increases.

8.2 Gene Regulatory Network Dynamics

To test whether the effect depends on oscillatory dynamics, we implemented a gene regulatory network (GRN) model with Hill-function activation:

$$\dot{x}_i = -\frac{x_i}{\tau} + \sigma \left(\sum_j W_{ij} x_j \right) + \eta_i(t), \quad (5)$$

where $\sigma(\cdot)$ is a sigmoid activation function, W is a sparse random connectivity matrix, and τ is the decay timescale.

The “code” in this context represents transcription factor readout: a low-dimensional projection (via DCT) of one network’s expression state constrains the other network’s dynamics.

Results (Figure 5, $N = 256$ genes) show the same qualitative signature as the Kuramoto system: $N_{\text{eff}}(A)$ remains constant at ≈ 193 , while $N_{\text{eff}}(B)$ collapses from 192 at $k = 1$ to 121 at $k = 32$ (37% reduction). Mismatch increases modestly (0.22 to 0.37) as bandwidth decreases.

The GRN results demonstrate that complexity collapse under bandwidth-limited coupling is not specific to phase oscillators. The same signature emerges in a fundamentally different dynamical system—one that more directly models gene regulatory interactions. This supports the claim that the constraint mechanism is general, not an artifact of Kuramoto dynamics.

8.3 Alternative Complexity Metrics

A potential concern is that measuring complexity via spectral entropy of Fourier amplitudes is tautological, since the code is also Fourier-based. We therefore computed three additional

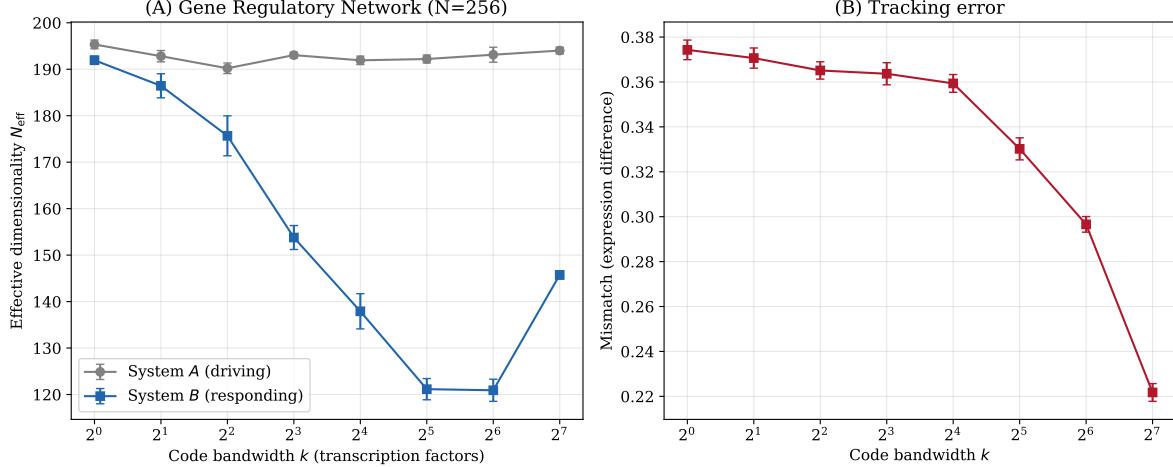


Figure 5: Complexity collapse in gene regulatory network dynamics ($N = 256$). (A) Effective dimensionality as a function of code bandwidth. System A (gray) maintains constant complexity while System B (blue) shows 37% reduction as k decreases. (B) Tracking error (expression difference) increases modestly at low bandwidth. The same constraint signature emerges in a fundamentally different dynamical system.

metrics that do not depend on the Fourier basis (Supplementary Figure S3):

- **PCA participation ratio:** effective dimensionality from covariance eigenvalues, $D_{\text{eff}} = (\sum \lambda)^2 / \sum \lambda^2$
- **Spatial gradient energy:** mean squared phase gradient $\langle |\nabla \theta|^2 \rangle$, measuring spatial smoothness
- **Kuramoto order parameter:** global synchrony $R = |\langle e^{i\theta} \rangle|$

All three metrics show the same qualitative pattern: System B 's organization changes systematically with code bandwidth (lower gradient energy and higher order at low k), while System A remains constant. This confirms that complexity collapse is a genuine dynamical phenomenon, not an artifact of the measurement basis.

8.4 Coupling Strength Robustness

We verified that the qualitative pattern persists across a range of coupling strengths $\lambda \in [0.25, 4.0]$ (Supplementary Figure S4). Stronger coupling produces greater complexity col-

lapse and tighter tracking, but the signature—complexity collapse with bounded mismatch—appears across the tested range.

9 Implications for Practice

The constraint interpretation suggests practical changes to how researchers across fields approach low-dimensional descriptions. We offer these not as prescriptions but as reframings that may prove useful.

Choosing reaction coordinates in molecular dynamics: The standard criterion for a good reaction coordinate is that it captures the slowest relaxation timescale (i.e., separates folded from unfolded). The constraint perspective suggests an additional criterion: does the coordinate *restrict accessible conformational dynamics* in the bound or functional state? A coordinate that merely summarizes the ensemble (thermometer) differs from one that actively constrains it (thermostat). The signature would be complexity collapse in conformational fluctuations when the coordinate is held fixed.

Interpreting single-cell manifolds: When UMAP or diffusion maps reveal low-dimensional structure, the standard interpretation is that this reflects underlying regulatory simplicity. The constraint perspective suggests an alternative: low-dimensional structure may reflect bandwidth limitations of inter-cellular coupling (e.g., signaling through a limited number of receptor types). Testing this would involve comparing manifold dimensionality in tissues with different coupling architectures.

Designing interfaces between models: Multi-scale modeling often couples coarse and fine scales through ad hoc “handshaking” variables. The constraint framework suggests a principled criterion: the interface should select coherent collective variables that induce complexity collapse in the fine-scale system without unbounded tracking error. Random or arbitrary coupling will fail this test.

Understanding why some low-D codes work: Neural population codes, morpho-

genetic gradients, and genetic regulatory motifs all achieve robust function through low-dimensional interfaces. The constraint interpretation explains *why* they work: they capture coherent collective variables that constrain downstream dynamics onto viable trajectories. Codes fail when they project onto incoherent modes that cannot stabilize the responding system.

Distinguishing constraint from information: Mutual information measures how much an observer can learn about one system from another. But a thermostat and a thermometer have the same mutual information with temperature—only one constrains it. When evaluating biological codes, ask: does reducing the code’s bandwidth produce complexity collapse (constraint) or tracking failure (information loss)? The two functions are distinct and should not be conflated.

10 Conclusion

This review has argued that the recurring interpretive problems surrounding dimensional mismatch in biology—sloppiness in protein modeling, artifacts in single-cell manifolds, shadows in neural codes—arise from a common source: the assumption that dimensional reduction preserves information rather than imposing constraints.

We proposed a unifying interpretation: low-dimensional interfaces between coupled high-dimensional systems function as stabilizing constraints rather than information channels. The signature of effective constraint is complexity collapse with bounded tracking: the responding system’s effective dimensionality decreases systematically with code bandwidth while maintaining alignment with the coarse-grained representation. Critically, this requires structured projections onto coherent collective variables; random projections fail.

The framework offers practical guidance across fields: criteria for choosing reaction coordinates in molecular dynamics, for interpreting manifold structure in single-cell data, for designing interfaces in multi-scale models, and for distinguishing constraint from informa-

tion in any coding system. The key question is not “how much information does the code preserve?” but “does the code constrain downstream dynamics onto viable trajectories?”

Beyond biology, this architecture suggests a blueprint for *neuromorphic constraint computing*, where system states are controlled not by precise error-correction of individual bits, but by shaping the topological manifolds of high-dimensional substrates.

The deeper insight may be that biological codes are not optimized for information transmission at all. They are dimensional valves that enable persistent organization by constraining what responding systems can do. Constraint and representation are distinct functions, and conflating them has systematically misled our understanding of biological organization.

Acknowledgments

The author thanks the reviewers for constructive feedback. No external funding was received for this work.

Declaration of generative AI use in the writing process

During the preparation of this work the author used Claude (Anthropic) for manuscript drafting, editing, and code development. The author reviewed and edited the content as needed and takes full responsibility for the content of the published article.

CRediT author statement

Ian Todd: Conceptualization, Methodology, Software, Formal analysis, Investigation, Writing – Original Draft, Writing – Review & Editing, Visualization.

Data availability

All simulation code and data are publicly available at <https://github.com/todd866/code-constraint>.

References

- D. Balduzzi and G. Tononi. Integrated information in discrete dynamical systems: motivation and theoretical framework. *PLoS Computational Biology*, 4(6):e1000091, 2008.
- R. B. Best and G. Hummer. Reaction coordinates and rates from transition paths. *Proceedings of the National Academy of Sciences*, 102(19):6732–6737, 2005.
- G. Buzsáki and A. Draguhn. Neuronal oscillations in cortical networks. *Science*, 304(5679):1926–1929, 2004.
- T. Chari and L. Pachter. The specious art of single-cell genomics. *PLoS Computational Biology*, 19(8):e1011288, 2023.
- S. M. Cooley, T. Hamilton, E. J. Deeds, and J. C. J. Ray. A novel metric reveals previously unrecognized distortion in dimensionality reduction of scRNA-seq data. *bioRxiv*, page 689851, 2019.
- J. P. Crutchfield and K. Young. Inferring statistical complexity. *Physical Review Letters*, 63(2):105–108, 1989.
- K. A. Dill and H. S. Chan. From Levinthal to pathways to funnels. *Nature Structural Biology*, 4(1):10–19, 1997.
- J. C. Flack. Coarse-graining as a downward causation mechanism. *Philosophical Transactions of the Royal Society A*, 375(2109):20160338, 2017.
- H. Frauenfelder, S. G. Sligar, and P. G. Wolynes. The energy landscapes and motions of proteins. *Science*, 254(5038):1598–1603, 1991.

- K. Friston. Life as we know it. *Journal of the Royal Society Interface*, 10(86):20130475, 2013.
- J. A. Gallego, M. G. Perich, L. E. Miller, and S. A. Solla. Neural manifolds for the control of movement. *Neuron*, 94(5):978–984, 2017.
- A. P. Georgopoulos, A. B. Schwartz, and R. E. Kettner. Neuronal population coding of movement direction. *Science*, 233(4771):1416–1419, 1986.
- R. N. Gutenkunst, J. J. Waterfall, F. P. Casey, K. S. Brown, C. R. Myers, and J. P. Sethna. Universally sloppy parameter sensitivities in systems biology models. *PLoS Computational Biology*, 3(10):e189, 2007.
- H. Haken. *Synergetics: An Introduction*. Springer, 1983.
- K. Henzler-Wildman and D. Kern. Dynamic personalities of proteins. *Nature*, 450(7172):964–972, 2007.
- D. Kobak and P. Berens. The art of using t-SNE for single-cell transcriptomics. *Nature Communications*, 10:5416, 2019.
- M. Levin. Molecular bioelectricity: how endogenous voltage potentials control cell behavior and instruct pattern regulation in vivo. *Molecular Biology of the Cell*, 25(24):3835–3850, 2014.
- M. Levin. Bioelectric signaling: Reprogrammable circuits underlying embryogenesis, regeneration, and cancer. *Cell*, 184(6):1971–1989, 2021.
- E. K. Miller, M. Lundqvist, and A. M. Bastos. Working memory 2.0. *Neuron*, 100(2):463–475, 2018.
- F. Noé and C. Clementi. Collective variables for the study of long-time kinetics from molecular trajectories: theory and methods. *Current Opinion in Structural Biology*, 43:141–147, 2017.

- V. S. Pande, K. Beauchamp, and G. R. Bowman. Everything you wanted to know about Markov State Models but were afraid to ask. *Methods*, 52(1):99–105, 2010.
- A. Pouget, P. Dayan, and R. Zemel. Information processing with population codes. *Nature Reviews Neuroscience*, 1(2):125–132, 2000.
- P. T. Sadtler, K. M. Quick, M. D. Golub, S. M. Chase, S. I. Ryu, E. C. Tyler-Kabara, B. M. Yu, and A. P. Batista. Neural constraints on learning. *Nature*, 512(7515):423–426, 2014.
- N. Tishby, F. C. Pereira, and W. Bialek. The information bottleneck method. In *Proceedings of the 37th Annual Allerton Conference on Communication, Control, and Computing*, pages 368–377, 1999.
- M. K. Transtrum, B. B. Machta, K. S. Brown, B. C. Daniels, C. R. Myers, and J. P. Sethna. Perspective: Sloppiness and emergent theories in physics, biology, and beyond. *The Journal of Chemical Physics*, 143(1):010901, 2015.
- C. Trapnell. Defining cell types and states with single-cell genomics. *Genome Research*, 25(10):1491–1498, 2015.
- A. M. Turing. The chemical basis of morphogenesis. *Philosophical Transactions of the Royal Society of London B*, 237(641):37–72, 1952.
- C. Weinreb, S. Wolock, B. K. Tusi, M. Socolovsky, and A. M. Klein. Fundamental limits on dynamic inference from single-cell snapshots. *Proceedings of the National Academy of Sciences*, 115(10):E2467–E2476, 2018.